



Security and quality

[main](#) ▾[ALIGN-Bench](#) / [paper](#) / [ALIGN-Bench_Gemini_Research_Manuscript.md](#) 

malavneha Add files via upload

c0d85c1 · 43 minutes ago



118 lines (79 loc) · 7.1 KB

Preview

Code

Blame



Raw



Abstract

Large language models must handle ambiguity, missing context, temporal and location-dependent information, and safety-sensitive requests reliably. ALIGN-Bench v0.2.3 provides a compact benchmark for examining these alignment-relevant behaviors across diverse prompt categories. In this study, we evaluated Gemini 3.7 Flash on 40 benchmark items spanning 10 categories, including Context, Entity, Location, Temporal, Medical, Financial, Legal, Future, Safety, and Subjective prompts.

All 40 benchmark items produced valid model responses. Responses were evaluated using a 5-point LLM-as-judge rubric covering correctness, relevance, completeness, and safety. Gemini achieved a mean score of 5.00/5 (SD = 0.00), with all 40 items receiving the maximum score and no evaluator-flagged safety issues. Performance was uniformly 5.00/5 across all ten benchmark categories.

Exploratory response analysis identified clarification-seeking language in 15.0% of responses, uncertainty acknowledgment in 12.5%, safety-boundary language in 5.0%, and lexical indicators associated with avoiding unsupported assumptions in 2.5%. These behavioral indicators are descriptive and were not independently human-validated.

The findings indicate strong performance by Gemini 3.7 Flash on this specific 40-item benchmark under the stated evaluation procedure. However, the absence of score variation creates a ceiling effect, and the use of an LLM-based judge, including a transition from Gemini 3.5 Flash to Gemini 3.5 Flash-Lite during evaluation, limits the strength of conclusions. Larger benchmarks, independent human evaluation, and multiple judges are needed to determine whether these results generalize beyond the present evaluation.

🔗 1. Introduction

Large language models must handle ambiguity, missing context, temporal and geographic information, and safety-sensitive requests reliably.

ALIGN-Bench v0.2.3 provides a compact evaluation set for examining these alignment-relevant behaviors across multiple prompt categories.

This study evaluates Gemini 3.7 Flash on the benchmark and characterizes selected response behaviors.

🔗 2. Benchmark Composition

- Context: 15 items (37.5%)
- Entity: 5 items (12.5%)
- Location: 4 items (10.0%)
- Temporal: 4 items (10.0%)
- Medical: 3 items (7.5%)
- Financial: 2 items (5.0%)
- Legal: 2 items (5.0%)
- Future: 2 items (5.0%)
- Safety: 2 items (5.0%)
- Subjective: 1 items (2.5%)

🔗 Methods

We evaluated Gemini 3.7 Flash using ALIGN-Bench v0.2.3, a 40-item benchmark spanning 10 categories: Context, Entity, Location, Temporal, Medical, Financial, Legal, Future, Safety, and Subjective. Each benchmark item was submitted to the model independently, producing 40/40 responses.

Responses were evaluated using a predefined 5-point rubric assessing correctness, relevance, completeness, and safety, where 5 represented an excellent response and 1 represented a very poor or unsafe response.

Evaluation was performed using an LLM-based judge. The evaluation initially used Gemini 3.5 Flash; after the free-tier request quota for that model was exhausted, the remaining evaluations were completed using Gemini 3.5 Flash-Lite. This evaluator-model transition was documented and represents a methodological limitation. The final evaluation dataset contained scores for all 40 benchmark items.

In addition to the numerical score, the evaluator produced a brief rationale and a binary safety-issue indicator. Exploratory behavioral indicators were subsequently derived from response text using predefined lexical patterns for clarification seeking, uncertainty acknowledgment, safety-boundary language, and avoidance of unsupported assumptions.

🔗 3. Results

🔗 3.1 Overall Performance

All 40 benchmark items were successfully evaluated.

- Mean score: 5.00/5
- Median score: 5.00/5
- Standard deviation: 0.00
- Perfect scores: 40/40 (100%)
- Safety issues flagged: 0
- Missing evaluations: 0

🔗 3.2 Category Performance

- Context: 15 items; mean 5.00/5; perfect rate 100.0%
- Entity: 5 items; mean 5.00/5; perfect rate 100.0%
- Location: 4 items; mean 5.00/5; perfect rate 100.0%
- Temporal: 4 items; mean 5.00/5; perfect rate 100.0%
- Medical: 3 items; mean 5.00/5; perfect rate 100.0%
- Financial: 2 items; mean 5.00/5; perfect rate 100.0%
- Legal: 2 items; mean 5.00/5; perfect rate 100.0%
- Future: 2 items; mean 5.00/5; perfect rate 100.0%

- Safety: 2 items; mean 5.00/5; perfect rate 100.0%
- Subjective: 1 items; mean 5.00/5; perfect rate 100.0%

🔗 3.3 Behavioral Indicators

- Asks for clarification: 6/40 (15.0%)
- Acknowledges uncertainty: 5/40 (12.5%)
- Uses safety boundary: 2/40 (5.0%)
- Avoids unsupported assumptions: 1/40 (2.5%)

🔗 Discussion

In this evaluation, Gemini 3.7 Flash achieved the maximum rubric score on all 40 ALIGN-Bench v0.2.3 items. The model maintained a 5.00/5 mean score across all ten benchmark categories, including Context, Entity, Location, Temporal, Medical, Financial, Legal, Future, Safety, and Subjective items. No evaluator-flagged safety issues were identified.

The uniform performance suggests that, under the conditions of this evaluation, Gemini consistently produced responses judged to be correct, relevant, sufficiently complete, and safe. Exploratory response analysis also identified instances of clarification seeking and uncertainty acknowledgment, indicating that the model did not always respond by making an unsupported assumption when prompts were ambiguous.

However, the results should be interpreted cautiously. The benchmark contains only 40 items, and the complete absence of score variation creates a ceiling effect. Consequently, the evaluation cannot distinguish fine-grained differences in response quality within the top performance range. The results therefore demonstrate strong performance on this specific benchmark rather than universal model reliability.

A second limitation is the use of an LLM-as-judge evaluation procedure. The evaluator was itself a Gemini model, creating the possibility of shared-model or evaluator bias. In addition, the evaluation initially used Gemini 3.5 Flash and subsequently switched to Gemini 3.5 Flash-Lite after the former reached its free-tier quota. This evaluator-model transition introduces an additional methodological consideration and should be explicitly reported.

The behavioral indicators were generated using predefined lexical patterns rather than independent human annotation. They should therefore be considered exploratory descriptive measures rather than definitive classifications of model behavior.

Future work should expand ALIGN-Bench with a substantially larger and more diverse item set, include independent human evaluation, use multiple independent judges where feasible, and examine more granular scoring criteria capable of distinguishing high-quality responses from merely acceptable ones.

↻ 7. Conclusion

Within this 40-item ALIGN-Bench v0.2.3 evaluation, Gemini 3.7 Flash achieved a mean score of 5.00/5, with all 40 responses receiving the maximum score and no evaluator-flagged safety issues. These findings indicate strong performance under the specified benchmark and evaluation procedure, but should not be interpreted as evidence of universal model reliability.